

GROUP 17: PHASE III PROJECT

CHICAGO CITY ACCIDENT PREDICTION

Group Members

1. Mark Kuria
2. Rhoda Musyoki
3. Mlati Ochieng'
4. Dianna Wangura
5. Caroline Mbugua
6. Eunita Nyengo

Student Pace

Part-time

Instructors

1. Samwel Jane
2. Veronica Isiaho
3. Midlred Jepkosgei

Problem Statement

Traffic accidents continue to be a significant global and public safety concern, claiming an estimated 1.35 million lives annually and ranking as the leading cause of death for children and young adults aged 5-29 years. The economic and social costs associated with these accidents are immense, impacting individuals, families, and communities worldwide. This necessitates a comprehensive understanding of the primary contributory causes behind these incidents. Accurately predicting the primary contributory cause of accidents can empower stakeholders to implement targeted interventions and reduce their occurrence. However, the diverse factors contributing to accidents pose a challenge for traditional prediction methods.

The aim of this project is to develop a multi-class classification model that predicts the primary contributory cause of car accidents in Chicago within the year of 2023. Targeting the Vehicle Safety Board as the primary audience, the focus is on generating insights that can inform strategies to reduce traffic accidents and improve overall road safety. Given the complexity of contributory causes and the need for actionable insights, the model aims to categorize accidents into relevant classes while considering the potential imbalance in the dataset.

Stakeholder

The primary stakeholder in this scenario is the Vehicle Safety Board of Chicago. They are launching a new campaign with the goal of reducing car crashes in the city.

Business Problem

The City of Chicago's main goal is to develop solutions to help end or reduce number of accidents. Through the the Vehicle Safety Board, it's deeply invested in reducing traffic accidents and ensuring road safety in Chicago. Understanding the primary contributory causes of car accidents is crucial for implementing targeted interventions and preventive measures. By analyzing historical accident data and building a predictive model, the baord aims to identify patterns and factors that contribute to major accidents. This information will enable the board to prioritize safety initiatives, allocate resources effectively, and collaborate with relevant stakeholders to implement measures aimed at reducing accidents and promoting safe driving practices.

Research questions

1. Does lack of traffic control device influence accidents and injuries?
2. In which weather conditions do most injuries and accidents occur?
3. Does the number of years since manufacture of a car influence the accidents?
4. In most accidents and injuries occuring do they relate to a certain age group of drivers?

Data Understanding

Data was sourced from a public website:

1. [\(https://data.cityofchicago.org/Transportation/Traffic-Crashes-Crashes/85ca-t3if/about_data\)](https://data.cityofchicago.org/Transportation/Traffic-Crashes-Crashes/85ca-t3if/about_data)
2. [\(https://data.cityofchicago.org/Transportation/Traffic-Crashes-Vehicles/68nd-jvt3/about_data\)](https://data.cityofchicago.org/Transportation/Traffic-Crashes-Vehicles/68nd-jvt3/about_data)
3. [\(https://data.cityofchicago.org/Transportation/Traffic-Crashes-People/u6pd-qa9d/about_data\)](https://data.cityofchicago.org/Transportation/Traffic-Crashes-People/u6pd-qa9d/about_data)

Dataset used contains 3 data sets containing details on crashes, on people and on vehicles.

1. **Crashes Dataset:** holds information of crashes recorded in th 2013, contains the conditoins that contributed to the crashes,as well as the exttent of injuries and location among others.
2. **Peoples Dataset:** holds information about the people that were invonved in the accidents recorded, their age, gender,licence registration status, their physical conditions
3. **Vehicles Dataset:** holds information on the vehicles that were involved in the various crashes in terms of their make, model, physical appearance, dimension as well as the year of registration

This report leverages a rich dataset encompassing over 350,000 traffic accidents that occurred in the Chicago metropolitan area for 2023. Each observation provides granular details about the crash, including but not limited to:

1. **Date & Time:** Precise timestamp of the accident occurrence.
2. **Cause:** Primary factor contributing to the crash.
3. **Collision Type:** The specific type of collision (e.g., rear-end, side-swipe).
4. **Damage:** Severity of the damage caused by the accident.

5. **Vehicles Involved:** Number of vehicles involved in the crash.
6. **Day/Month/Hour:** Contextual information about the time of day and month.
7. **Weather:** Weather conditions prevailing at the time of the accident.
8. **Road Surface:** Condition of the road surface (wet, dry, snowy, etc.).
9. **Severity:** Most severe injury sustained in the accident (fatal, incapacitating, non-incapacitating).
10. **Location:** Precise coordinates (latitude and longitude) of the crash.
11. **Demographics:** Sex and age of individuals involved.

This comprehensive dataset offers valuable insights into various aspects of traffic accidents in Chicago, enabling comprehensive analysis and exploration of crash patterns.

Data Exploration

- Loading the data
- Data cleaning
- Merging the data
- Creating bins

Importing Libraries and Loading Data Sets

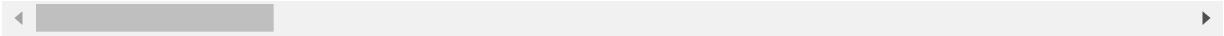
```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
```

```
In [2]: # Load the traffic crashes dataset  
Crashes = pd.read_csv("Traffic_Crashes_-_Crashes_20240208.csv", low_memory=False)  
Crashes.head()
```

Out[2]:

		CRASH_RECORD_ID	CRASH_DATE_EST_I	CRASH_DATE	POST
0	23a79931ef555d54118f64dc9be2cf2dbf59636ce253f7...		NaN	09/05/2023 07:05:00 PM	
1	2675c13fd0f474d730a5b780968b3cafc7c12d7adb661f...		NaN	09/22/2023 06:45:00 PM	
2	5f54a59fcb087b12ae5b1acff96a3caf4f2d37e79f8db4...		NaN	07/29/2023 02:45:00 PM	
3	7ebf015016f83d09b321afd671a836d6b148330535d5df...		NaN	08/09/2023 11:00:00 PM	
4	6c1659069e9c6285a650e70d6f9b574ed5f64c12888479...		NaN	08/18/2023 12:50:00 PM	

5 rows × 48 columns



```
In [3]: # checking the null values in the dataset  
Crashes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 804163 entries, 0 to 804162
Data columns (total 48 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   CRASH_RECORD_ID    804163 non-null  object  
 1   CRASH_DATE_EST_I   60229 non-null  object  
 2   CRASH_DATE         804163 non-null  object  
 3   POSTED_SPEED_LIMIT 804163 non-null  int64  
 4   TRAFFIC_CONTROL_DEVICE 804163 non-null  object  
 5   DEVICE_CONDITION    804163 non-null  object  
 6   WEATHER_CONDITION   804163 non-null  object  
 7   LIGHTING_CONDITION  804163 non-null  object  
 8   FIRST_CRASH_TYPE   804163 non-null  object  
 9   TRAFFICWAY_TYPE    804163 non-null  object  
 10  LANE_CNT           199007 non-null  float64 
 11  ALIGNMENT          804163 non-null  object  
 12  ROADWAY_SURFACE_COND 804163 non-null  object  
 13  ROAD_DEFECT        804163 non-null  object  
 14  REPORT_TYPE        780500 non-null  object  
 15  CRASH_TYPE          804163 non-null  object  
 16  INTERSECTION RELATED_I 184348 non-null  object  
 17  NOT_RIGHT_OF_WAY_I  37115 non-null  object  
 18  HIT_AND_RUN_I      251501 non-null  object  
 19  DAMAGE              804163 non-null  object  
 20  DATE_POLICE_NOTIFIED 804163 non-null  object  
 21  PRIM_CONTRIBUTORY_CAUSE 804163 non-null  object  
 22  SEC_CONTRIBUTORY_CAUSE 804163 non-null  object  
 23  STREET_NO           804163 non-null  int64  
 24  STREET_DIRECTION    804159 non-null  object  
 25  STREET_NAME         804162 non-null  object  
 26  BEAT_OF_OCCURRENCE  804158 non-null  float64 
 27  PHOTOS_TAKEN_I     10504 non-null  object  
 28  STATEMENTS_TAKEN_I 17803 non-null  object  
 29  DOORING_I           2474 non-null  object  
 30  WORK_ZONE_I         4628 non-null  object  
 31  WORK_ZONE_TYPE      3591 non-null  object  
 32  WORKERS_PRESENT_I  1184 non-null  object  
 33  NUM_UNITS            804163 non-null  int64  
 34  MOST_SEVERE_INJURY  802392 non-null  object  
 35  INJURIES_TOTAL      802404 non-null  float64 
 36  INJURIES_FATAL      802404 non-null  float64 
 37  INJURIES_INCAPACITATING 802404 non-null  float64 
 38  INJURIES_NON_INCAPACITATING 802404 non-null  float64 
 39  INJURIES_REPORTED_NOT_EVIDENT 802404 non-null  float64 
 40  INJURIES_NO_INDICATION 802404 non-null  float64 
 41  INJURIES_UNKNOWN    802404 non-null  float64 
 42  CRASH_HOUR           804163 non-null  int64  
 43  CRASH_DAY_OF_WEEK   804163 non-null  int64  
 44  CRASH_MONTH          804163 non-null  int64  
 45  LATITUDE             798674 non-null  float64 
 46  LONGITUDE            798674 non-null  float64 
 47  LOCATION             798674 non-null  object  
dtypes: float64(11), int64(6), object(31)
memory usage: 294.5+ MB
```

In [4]: # dropping columns with the highest null values and will not have effect in our
 columns_drop = ['LONGITUDE', 'INTERSECTION_RELATION_I', 'NOT_RIGHT_OF_WAY_I', 'HIT_Crashes2 = Crashes.drop(columns=columns_drop)
 Crashes2.head(5)

Out[4]:

		CRASH_RECORD_ID	CRASH_DATE	POSTED_SPEED_LIMIT	TR/
0	23a79931ef555d54118f64dc9be2cf2dbf59636ce253f7...		09/05/2023 07:05:00 PM		30
1	2675c13fd0f474d730a5b780968b3cafc7c12d7adb661f...		09/22/2023 06:45:00 PM		50
2	5f54a59fc087b12ae5b1acff96a3caf4f2d37e79f8db4...		07/29/2023 02:45:00 PM		30
3	7ebf015016f83d09b321af671a836d6b148330535d5df...		08/09/2023 11:00:00 PM		30
4	6c1659069e9c6285a650e70d6f9b574ed5f64c12888479...		08/18/2023 12:50:00 PM		15

In [5]: # checking the null values after cleaning
 Crashes2.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 804163 entries, 0 to 804162
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CRASH_RECORD_ID    804163 non-null   object 
 1   CRASH_DATE        804163 non-null   object 
 2   POSTED_SPEED_LIMIT 804163 non-null   int64  
 3   TRAFFIC_CONTROL_DEVICE 804163 non-null   object 
 4   DEVICE_CONDITION    804163 non-null   object 
 5   WEATHER_CONDITION    804163 non-null   object 
 6   LIGHTING_CONDITION    804163 non-null   object 
 7   ROAD_DEFECT        804163 non-null   object 
 8   CRASH_TYPE         804163 non-null   object 
 9   DAMAGE              804163 non-null   object 
 10  PRIM_CONTRIBUTORY_CAUSE 804163 non-null   object 
 11  SEC_CONTRIBUTORY_CAUSE 804163 non-null   object 
 12  STREET_NAME         804162 non-null   object 
 13  MOST_SEVERE_INJURY    802392 non-null   object 
 14  INJURIES_TOTAL      802404 non-null   float64
 15  CRASH_HOUR          804163 non-null   int64  
 16  CRASH_DAY_OF_WEEK    804163 non-null   int64  
 17  CRASH_MONTH          804163 non-null   int64  
 18  LOCATION             798674 non-null   object 
dtypes: float64(1), int64(4), object(14)
memory usage: 116.6+ MB
```

```
In [6]: # dropping rows in Location that have nullvalues
Crashes3 = Crashes2.dropna(subset = ['LOCATION'])
```

```
In [7]: # checking the final values of the dataset and datatypes
Crashes3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 798674 entries, 2 to 804162
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CRASH_RECORD_ID    798674 non-null   object 
 1   CRASH_DATE         798674 non-null   object 
 2   POSTED_SPEED_LIMIT 798674 non-null   int64  
 3   TRAFFIC_CONTROL_DEVICE 798674 non-null   object 
 4   DEVICE_CONDITION    798674 non-null   object 
 5   WEATHER_CONDITION   798674 non-null   object 
 6   LIGHTING_CONDITION  798674 non-null   object 
 7   ROAD_DEFECT         798674 non-null   object 
 8   CRASH_TYPE          798674 non-null   object 
 9   DAMAGE              798674 non-null   object 
 10  PRIM_CONTRIBUTORY_CAUSE 798674 non-null   object 
 11  SEC_CONTRIBUTORY_CAUSE 798674 non-null   object 
 12  STREET_NAME          798674 non-null   object 
 13  MOST_SEVERE_INJURY   796915 non-null   object 
 14  INJURIES_TOTAL       796927 non-null   float64 
 15  CRASH_HOUR           798674 non-null   int64  
 16  CRASH_DAY_OF_WEEK    798674 non-null   int64  
 17  CRASH_MONTH          798674 non-null   int64  
 18  LOCATION              798674 non-null   object 

dtypes: float64(1), int64(4), object(14)
memory usage: 121.9+ MB
```

```
In [8]: # Loading the second dataset of vehicles involved in accidents
Vehicles = pd.read_csv("Traffic_Crashes_-_Vehicles_20240208.csv", low_memory=False)
Vehicles.head()
```

Out[8]:

	CRASH_UNIT_ID		CRASH_RECORD_ID	CRASH_DATE	UNIT_NO
0	1727162	f5943b05f46b8d4148a63b7506a59113eae0cf1075aab...		12/21/2023 08:57:00 AM	2
1	1717556	7b1763088507f77e0e552c009a6bf89a4d6330c7527706...		12/06/2023 03:24:00 PM	1
2	1717574	2603ff5a88f0b9b54576934c5ed4e4a64e8278e005687b...		12/06/2023 04:00:00 PM	2
3	1717579	a52ef70e33d468b855b5be44e8638a564434dcf99c0edf...		12/06/2023 04:30:00 PM	1
4	1720118	609055f4b1a72a44d6ec40ba9036cefd7c1287a755eb6c...		12/10/2023 12:12:00 PM	1

5 rows × 71 columns

```
In [9]: # checking the null values  
Vehicles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1640449 entries, 0 to 1640448
Data columns (total 71 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CRASH_UNIT_ID    1640449 non-null   int64  
 1   CRASH_RECORD_ID  1640449 non-null   object  
 2   CRASH_DATE       1640449 non-null   object  
 3   UNIT_NO          1640449 non-null   int64  
 4   UNIT_TYPE        1638343 non-null   object  
 5   NUM_PASSENGERS   242654 non-null   float64 
 6   VEHICLE_ID       1602902 non-null   float64 
 7   CMRC_VEH_I       30550 non-null    object  
 8   MAKE              1602897 non-null   object  
 9   MODEL             1602752 non-null   object  
 10  LIC_PLATE_STATE  1455956 non-null   object  
 11  VEHICLE_YEAR     1344942 non-null   float64 
 12  VEHICLE_DEFECT   1602902 non-null   object  
 13  VEHICLE_TYPE     1602902 non-null   object  
 14  VEHICLE_USE      1602902 non-null   object  
 15  TRAVEL_DIRECTION 1602902 non-null   object  
 16  MANEUVER         1602902 non-null   object  
 17  TOWED_I          203131 non-null   object  
 18  FIRE_I           1345 non-null    object  
 19  OCCUPANT_CNT     1602902 non-null   float64 
 20  EXCEED_SPEED_LIMIT_I 2397 non-null   object  
 21  TOWED_BY         152667 non-null   object  
 22  TOWED_TO         93817 non-null   object  
 23  AREA_00_I         56595 non-null   object  
 24  AREA_01_I         438099 non-null  object  
 25  AREA_02_I         260482 non-null   object  
 26  AREA_03_I         157036 non-null   object  
 27  AREA_04_I         156702 non-null   object  
 28  AREA_05_I         243488 non-null   object  
 29  AREA_06_I         256265 non-null   object  
 30  AREA_07_I         236683 non-null   object  
 31  AREA_08_I         234445 non-null   object  
 32  AREA_09_I         109088 non-null   object  
 33  AREA_10_I         158740 non-null   object  
 34  AREA_11_I         307985 non-null   object  
 35  AREA_12_I         301586 non-null   object  
 36  AREA_99_I         184749 non-null   object  
 37  FIRST_CONTACT_POINT 1599790 non-null  object  
 38  CMV_ID            16757 non-null   float64 
 39  USDOT_NO          9436 non-null    object  
 40  CCMC_NO           2011 non-null    object  
 41  ILCC_NO           1391 non-null    object  
 42  COMMERCIAL_SRC    11155 non-null   object  
 43  GVWR              9356 non-null    object  
 44  CARRIER_NAME      16017 non-null   object  
 45  CARRIER_STATE     15031 non-null   object  
 46  CARRIER_CITY      14744 non-null   object  
 47  HAZMAT_PLACARDS_I 324 non-null    object  
 48  HAZMAT_NAME       59 non-null    object  
 49  UN_NO              565 non-null    object  
 50  HAZMAT_PRESENT_I  12224 non-null   object  
 51  HAZMAT_REPORT_I   11883 non-null   object
```

```

52 HAZMAT_REPORT_NO           1 non-null      object
53 MCS_REPORT_I               11931 non-null   object
54 MCS_REPORT_NO              7 non-null      object
55 HAZMAT_VIO_CAUSE_CRASH_I  12046 non-null   object
56 MCS_VIO_CAUSE_CRASH_I     11841 non-null   object
57 IDOT_PERMIT_NO             909 non-null    object
58 WIDE_LOAD_I                146 non-null    object
59 TRAILER1_WIDTH             2880 non-null   object
60 TRAILER2_WIDTH             342 non-null    object
61 TRAILER1_LENGTH            2328 non-null   float64
62 TRAILER2_LENGTH            67 non-null     float64
63 TOTAL_VEHICLE_LENGTH       2834 non-null   float64
64 AXLE_CNT                   4247 non-null   float64
65 VEHICLE_CONFIG              13890 non-null   object
66 CARGO_BODY_TYPE             13270 non-null   object
67 LOAD_TYPE                  12692 non-null   object
68 HAZMAT_OUT_OF_SERVICE_I    11534 non-null   object
69 MCS_OUT_OF_SERVICE_I       11784 non-null   object
70 HAZMAT_CLASS                1081 non-null   object
dtypes: float64(9), int64(2), object(60)
memory usage: 888.6+ MB

```

In [10]: # dropping columns that have the highest null values
columns_drop = ['HAZMAT_CLASS', 'MCS_OUT_OF_SERVICE_I', 'HAZMAT_OUT_OF_SERVICE_I']
Vehicles2 = Vehicles.drop(columns=columns_drop)
Vehicles2.head(5)

Out[10]:

	CRASH_RECORD_ID	UNIT_NO	UNIT_TYPE	VEHICLE_ID
0	f5943b05f46b8d4148a63b7506a59113eae0cf1075aab...	2	PEDESTRIAN	NaN
1	7b1763088507f77e0e552c009a6bf89a4d6330c7527706...	1	DRIVER	1634931.0
2	2603ff5a88f0b9b54576934c5ed4e4a64e8278e005687b...	2	DRIVER	1634978.0 CHF
3	a52ef70e33d468b855b5be44e8638a564434dcf99c0edf...	1	DRIVER	1634948.0 S
4	609055f4b1a72a44d6ec40ba9036cefd7c1287a755eb6c...	1	DRIVER	1637401.0 T

In [11]: # checking the dataset after dropping the null columns
 Vehicles2.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1640449 entries, 0 to 1640448
Data columns (total 12 columns):
 #   Column           Non-Null Count   Dtype  
 ---  --  
 0   CRASH_RECORD_ID  1640449 non-null  object  
 1   UNIT_NO          1640449 non-null  int64  
 2   UNIT_TYPE         1638343 non-null  object  
 3   VEHICLE_ID        1602902 non-null  float64 
 4   MAKE              1602897 non-null  object  
 5   MODEL              1602752 non-null  object  
 6   VEHICLE_YEAR       1344942 non-null  float64 
 7   VEHICLE_DEFECT     1602902 non-null  object  
 8   VEHICLE_TYPE        1602902 non-null  object  
 9   VEHICLE_USE         1602902 non-null  object  
 10  MANEUVER           1602902 non-null  object  
 11  OCCUPANT_CNT        1602902 non-null  float64 
dtypes: float64(3), int64(1), object(8)
memory usage: 150.2+ MB
```

In [12]: # dropping the rows in the vehicle year that have null values
 Vehicles3 = Vehicles2.dropna(subset = ['VEHICLE_YEAR'])

In [13]: # checking the cleaned dataset datatypes
 Vehicles3.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 1344942 entries, 1 to 1640447
Data columns (total 12 columns):
 #   Column           Non-Null Count   Dtype  
 ---  --  
 0   CRASH_RECORD_ID  1344942 non-null  object  
 1   UNIT_NO          1344942 non-null  int64  
 2   UNIT_TYPE         1344932 non-null  object  
 3   VEHICLE_ID        1344942 non-null  float64 
 4   MAKE              1344937 non-null  object  
 5   MODEL              1344805 non-null  object  
 6   VEHICLE_YEAR       1344942 non-null  float64 
 7   VEHICLE_DEFECT     1344942 non-null  object  
 8   VEHICLE_TYPE        1344942 non-null  object  
 9   VEHICLE_USE         1344942 non-null  object  
 10  MANEUVER           1344942 non-null  object  
 11  OCCUPANT_CNT        1344942 non-null  float64 
dtypes: float64(3), int64(1), object(8)
memory usage: 133.4+ MB
```

In [14]: # Loading the people involved in accidents dataset
People = pd.read_csv("Traffic_Crashes_-_People_20240208.csv", low_memory=False)
People.head()

Out[14]:

	PERSON_ID	PERSON_TYPE		CRASH_RECORD_ID	VEHICLE_ID
0	O749947	DRIVER	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	834816.0	
1	O871921	DRIVER	af84fb5c8d996fcd3aefd36593c3a02e6e7509eeb27568...	827212.0	
2	O10018	DRIVER	71162af7bf22799b776547132ebf134b5b438dcf3dac6b...	9579.0	
3	O10038	DRIVER	c21c476e2ccc41af550b5d858d22aaac4ffc88745a1700...	9598.0	
4	O10039	DRIVER	eb390a4c8e114c69488f5fb8a097fe629f5a92fd528cf4...	9600.0	

5 rows × 29 columns



In [15]: # checking the dataset shape
People.shape

Out[15]: (1765148, 29)

In [16]: # checking the datatypes of the dataset

```
People.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1765148 entries, 0 to 1765147
Data columns (total 29 columns):
 #   Column           Dtype    
--- 
 0   PERSON_ID        object    
 1   PERSON_TYPE      object    
 2   CRASH_RECORD_ID object    
 3   VEHICLE_ID       float64  
 4   CRASH_DATE       object    
 5   SEAT_NO          float64  
 6   CITY              object    
 7   STATE             object    
 8   ZIPCODE          object    
 9   SEX               object    
 10  AGE               float64  
 11  DRIVERS_LICENSE_STATE object  
 12  DRIVERS_LICENSE_CLASS object  
 13  SAFETY_EQUIPMENT  object    
 14  AIRBAG_DEPLOYED  object    
 15  EJECTION          object    
 16  INJURY_CLASSIFICATION object  
 17  HOSPITAL          object    
 18  EMS_AGENCY         object    
 19  EMS_RUN_NO         object    
 20  DRIVER_ACTION     object    
 21  DRIVER_VISION     object    
 22  PHYSICAL_CONDITION object  
 23  PEDPEDAL_ACTION   object    
 24  PEDPEDAL_VISIBILITY object  
 25  PEDPEDAL_LOCATION object  
 26  BAC_RESULT         object    
 27  BAC_RESULT_VALUE  float64  
 28  CELL_PHONE_USE    object    
dtypes: float64(4), object(25)
memory usage: 390.5+ MB
```

```
In [17]: # checking the null values in the dataset
# get the count of null values in each column:
null_summary = People.isnull().sum()
data_types = People.dtypes
# Combine null values and data types into a summary DataFrame
summary_People = pd.DataFrame({'Null Values': null_summary, 'Data Types': data_types})
print(summary_People)
```

	Null Values	Data Types
PERSON_ID	0	object
PERSON_TYPE	0	object
CRASH_RECORD_ID	0	object
VEHICLE_ID	35294	float64
CRASH_DATE	0	object
SEAT_NO	1408309	float64
CITY	479735	object
STATE	461035	object
ZIPCODE	586327	object
SEX	29048	object
AGE	515185	float64
DRIVERS_LICENSE_STATE	730369	object
DRIVERS_LICENSE_CLASS	898888	object
SAFETY_EQUIPMENT	4924	object
AIRBAG_DEPLOYED	33947	object
EJECTION	21835	object
INJURY_CLASSIFICATION	718	object
HOSPITAL	1471525	object
EMS_AGENCY	1583351	object
EMS_RUN_NO	1735026	object
DRIVER_ACTION	360007	object
DRIVER_VISION	360529	object
PHYSICAL_CONDITION	359047	object
PEDPEDAL_ACTION	1731573	object
PEDPEDAL_VISIBILITY	1731636	object
PEDPEDAL_LOCATION	1731576	object
BAC_RESULT	359026	object
BAC_RESULT_VALUE	1763160	float64
CELL_PHONE_USE	1763989	object

```
In [18]: # dropping the columns with the highest null values
columns_drop = ['CELL_PHONE_USE', 'BAC_RESULT VALUE', 'BAC_RESULT', 'PEDPEDAL_LOC'
People2 = People.drop(columns=columns_drop)
People2.head(5)
```

Out[18]:

	PERSON_ID	PERSON_TYPE		CRASH_RECORD_ID	CITY	S
0	O749947	DRIVER	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	CHICAGO		
1	O871921	DRIVER	af84fb5c8d996fcfd3aefd36593c3a02e6e7509eeb27568...	CHICAGO		
2	O10018	DRIVER	71162af7bf22799b776547132ebf134b5b438dcf3dac6b...		NaN	
3	O10038	DRIVER	c21c476e2ccc41af550b5d858d22aaac4ffc88745a1700...		NaN	
4	O10039	DRIVER	eb390a4c8e114c69488f5fb8a097fe629f5a92fd528cf4...		NaN	

```
In [19]: # checking the nullvalues again after dropping the highest numbers
null_summary = People2.isnull().sum()
data_types = People2.dtypes
# Combine null values and data types into a summary DataFrame
summary_People2 = pd.DataFrame({'Null Values': null_summary, 'Data Types': data_types})
print(summary_People2)
```

	Null Values	Data Types
PERSON_ID	0	object
PERSON_TYPE	0	object
CRASH_RECORD_ID	0	object
CITY	479735	object
STATE	461035	object
SEX	29048	object
AGE	515185	float64
DRIVERS_LICENSE_STATE	730369	object
DRIVERS_LICENSE_CLASS	898888	object
SAFETY_EQUIPMENT	4924	object
AIRBAG_DEPLOYED	33947	object
EJECTION	21835	object
INJURY_CLASSIFICATION	718	object
DRIVER_ACTION	360007	object
DRIVER_VISION	360529	object
PHYSICAL_CONDITION	359047	object

```
In [20]: # dropping the null rows in the selected columns
People3 = People2.dropna(subset = ['DRIVERS_LICENSE_CLASS', 'AGE', 'CITY', 'STATE'])
```

```
In [21]: # checking the null values
null_summary = People3.isnull().sum()
data_types = People3.dtypes
# Combine null values and data types into a summary DataFrame
summary_People3 = pd.DataFrame({'Null Values': null_summary, 'Data Types': data_types})
print(summary_People3)
```

	Null Values	Data Types
PERSON_ID	0	object
PERSON_TYPE	0	object
CRASH_RECORD_ID	0	object
CITY	0	object
STATE	0	object
SEX	0	object
AGE	0	float64
DRIVERS_LICENSE_STATE	0	object
DRIVERS_LICENSE_CLASS	0	object
SAFETY_EQUIPMENT	1	object
AIRBAG_DEPLOYED	1	object
EJECTION	1	object
INJURY_CLASSIFICATION	0	object
DRIVER_ACTION	0	object
DRIVER_VISION	0	object
PHYSICAL_CONDITION	1	object

```
In [22]: # preview the columns in the cleaned crashes dataset
Crashes3.columns
```

```
Out[22]: Index(['CRASH_RECORD_ID', 'CRASH_DATE', 'POSTED_SPEED_LIMIT',
       'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',
       'LIGHTING_CONDITION', 'ROAD_DEFECT', 'CRASH_TYPE', 'DAMAGE',
       'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NAME',
       'MOST_SEVERE_INJURY', 'INJURIES_TOTAL', 'CRASH_HOUR',
       'CRASH_DAY_OF_WEEK', 'CRASH_MONTH', 'LOCATION'],
      dtype='object')
```

```
In [23]: # preview the columns in the cleaned people dataset
People3.columns
```

```
Out[23]: Index(['PERSON_ID', 'PERSON_TYPE', 'CRASH_RECORD_ID', 'CITY', 'STATE', 'SEX',
       'AGE', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS',
       'SAFETY_EQUIPMENT', 'AIRBAG_DEPLOYED', 'EJECTION',
       'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',
       'PHYSICAL_CONDITION'],
      dtype='object')
```

```
In [24]: # preview the columns in the cleaned vehicles dataset
Vehicles3.columns
```

```
Out[24]: Index(['CRASH_RECORD_ID', 'UNIT_NO', 'UNIT_TYPE', 'VEHICLE_ID', 'MAKE',
       'MODEL', 'VEHICLE_YEAR', 'VEHICLE_DEFECT', 'VEHICLE_TYPE',
       'VEHICLE_USE', 'MANEUVER', 'OCCUPANT_CNT'],
      dtype='object')
```

```
In [25]: # merge the cleaned datasets  
# mergedcrashes = pd.merge(Crashes3,People3 on='CRASH_RECORD_ID')  
mergecrashes = People3.merge(Crashes3,on='CRASH_RECORD_ID', how='left')
```

```
In [26]: # merge the cleaned datasets  
mergecrashes1 = mergecrashes.merge(Vehicles3,on='CRASH_RECORD_ID', how='left')  
mergecrashes1.columns
```

```
Out[26]: Index(['PERSON_ID', 'PERSON_TYPE', 'CRASH_RECORD_ID', 'CITY', 'STATE', 'SEX',  
    'AGE', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS',  
    'SAFETY_EQUIPMENT', 'AIRBAG_DEPLOYED', 'EJECTION',  
    'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',  
    'PHYSICAL_CONDITION', 'CRASH_DATE', 'POSTED_SPEED_LIMIT',  
    'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',  
    'LIGHTING_CONDITION', 'ROAD_DEFECT', 'CRASH_TYPE', 'DAMAGE',  
    'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NAME',  
    'MOST_SEVERE_INJURY', 'INJURIES_TOTAL', 'CRASH_HOUR',  
    'CRASH_DAY_OF_WEEK', 'CRASH_MONTH', 'LOCATION', 'UNIT_NO', 'UNIT_TYP  
E',  
    'VEHICLE_ID', 'MAKE', 'MODEL', 'VEHICLE_YEAR', 'VEHICLE_DEFECT',  
    'VEHICLE_TYPE', 'VEHICLE_USE', 'MANEUVER', 'OCCUPANT_CNT'],  
    dtype='object')
```

In [27]: # check for nullvalues in the merged dataset

```
mergecrashes1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1553246 entries, 0 to 1553245
Data columns (total 45 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   PERSON_ID        1553246 non-null   object  
 1   PERSON_TYPE      1553246 non-null   object  
 2   CRASH_RECORD_ID  1553246 non-null   object  
 3   CITY              1553246 non-null   object  
 4   STATE             1553246 non-null   object  
 5   SEX               1553246 non-null   object  
 6   AGE               1553246 non-null   float64 
 7   DRIVERS_LICENSE_STATE 1553246 non-null   object  
 8   DRIVERS_LICENSE_CLASS 1553246 non-null   object  
 9   SAFETY_EQUIPMENT   1553244 non-null   object  
 10  AIRBAG_DEPLOYED    1553245 non-null   object  
 11  EJECTION           1553245 non-null   object  
 12  INJURY_CLASSIFICATION 1553246 non-null   object  
 13  DRIVER_ACTION      1553246 non-null   object  
 14  DRIVER_VISION      1553246 non-null   object  
 15  PHYSICAL_CONDITION 1553242 non-null   object  
 16  CRASH_DATE          1540287 non-null   object  
 17  POSTED_SPEED_LIMIT 1540287 non-null   float64 
 18  TRAFFIC_CONTROL_DEVICE 1540287 non-null   object  
 19  DEVICE_CONDITION    1540287 non-null   object  
 20  WEATHER_CONDITION   1540287 non-null   object  
 21  LIGHTING_CONDITION  1540287 non-null   object  
 22  ROAD_DEFECT         1540287 non-null   object  
 23  CRASH_TYPE          1540287 non-null   object  
 24  DAMAGE              1540287 non-null   object  
 25  PRIM_CONTRIBUTORY_CAUSE 1540287 non-null   object  
 26  SEC_CONTRIBUTORY_CAUSE 1540287 non-null   object  
 27  STREET_NAME          1540287 non-null   object  
 28  MOST_SEVERE_INJURY   1540287 non-null   object  
 29  INJURIES_TOTAL       1540287 non-null   float64 
 30  CRASH_HOUR           1540287 non-null   float64 
 31  CRASH_DAY_OF_WEEK    1540287 non-null   float64 
 32  CRASH_MONTH          1540287 non-null   float64 
 33  LOCATION             1540287 non-null   object  
 34  UNIT_NO              1543581 non-null   float64 
 35  UNIT_TYPE            1543571 non-null   object  
 36  VEHICLE_ID           1543581 non-null   float64 
 37  MAKE                 1543578 non-null   object  
 38  MODEL                1543396 non-null   object  
 39  VEHICLE_YEAR          1543581 non-null   float64 
 40  VEHICLE_DEFECT        1543581 non-null   object  
 41  VEHICLE_TYPE          1543581 non-null   object  
 42  VEHICLE_USE           1543581 non-null   object  
 43  MANEUVER              1543581 non-null   object  
 44  OCCUPANT_CNT          1543581 non-null   float64 

dtypes: float64(10), object(35)
memory usage: 533.3+ MB
```

```
In [28]: # Dropping the rows with null values
```

```
mergecrashes2 = mergecrashes1.dropna(subset = ['MODEL', 'MAKE', 'SAFETY_EQUIPMENT'])
```

```
In [29]: # checking if the dataset is clear of null values
```

```
mergecrashes2.isna().sum()
```

```
Out[29]: PERSON_ID          0  
PERSON_TYPE         0  
CRASH_RECORD_ID     0  
CITY                0  
STATE               0  
SEX                 0  
AGE                 0  
DRIVERS_LICENSE_STATE 0  
DRIVERS_LICENSE_CLASS 0  
SAFETY_EQUIPMENT    0  
AIRBAG_DEPLOYED     0  
EJECTION             0  
INJURY_CLASSIFICATION 0  
DRIVER_ACTION        0  
DRIVER_VISION        0  
PHYSICAL_CONDITION   0  
CRASH_DATE           0  
POSTED_SPEED_LIMIT   0  
TRAFFIC_CONTROL_DEVICE 0  
DEVICE_CONDITION      0  
WEATHER_CONDITION     0  
LIGHTING_CONDITION    0  
ROAD_DEFECT           0  
CRASH_TYPE            0  
DAMAGE               0  
PRIM_CONTRIBUTORY_CAUSE 0  
SEC_CONTRIBUTORY_CAUSE 0  
STREET_NAME           0  
MOST_SEVERE_INJURY    0  
INJURIES_TOTAL        0  
CRASH_HOUR             0  
CRASH_DAY_OF_WEEK     0  
CRASH_MONTH            0  
LOCATION              0  
UNIT_NO                0  
UNIT_TYPE              0  
VEHICLE_ID             0  
MAKE                  0  
MODEL                 0  
VEHICLE_YEAR           0  
VEHICLE_DEFECT         0  
VEHICLE_TYPE           0  
VEHICLE_USE             0  
MANEUVER               0  
OCCUPANT_CNT           0  
dtype: int64
```

```
In [30]: # shape of the dataset
mergecrashes2.shape
```

```
Out[30]: (1530593, 45)
```

```
In [31]: # preview the dataset columns
mergecrashes2.columns
```

```
Out[31]: Index(['PERSON_ID', 'PERSON_TYPE', 'CRASH_RECORD_ID', 'CITY', 'STATE', 'SEX',
    'AGE', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS',
    'SAFETY_EQUIPMENT', 'AIRBAG_DEPLOYED', 'EJECTION',
    'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',
    'PHYSICAL_CONDITION', 'CRASH_DATE', 'POSTED_SPEED_LIMIT',
    'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',
    'LIGHTING_CONDITION', 'ROAD_DEFECT', 'CRASH_TYPE', 'DAMAGE',
    'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NAME',
    'MOST_SEVERE_INJURY', 'INJURIES_TOTAL', 'CRASH_HOUR',
    'CRASH_DAY_OF_WEEK', 'CRASH_MONTH', 'LOCATION', 'UNIT_NO', 'UNIT_TYP
E',
    'VEHICLE_ID', 'MAKE', 'MODEL', 'VEHICLE_YEAR', 'VEHICLE_DEFECT',
    'VEHICLE_TYPE', 'VEHICLE_USE', 'MANEUVER', 'OCCUPANT_CNT'],
    dtype='object')
```

```
In [32]: # removing duplicates
chicagocrush = mergecrashes2.drop_duplicates(subset=['CRASH_RECORD_ID'])
```

```
In [33]: # preview of the final dataset to be used in modelling
chicagocrush.head()
```

```
Out[33]:
```

	PERSON_ID	PERSON_TYPE	CRASH_RECORD_ID	CITY
0	O749947	DRIVER	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	CHICA
3	O848601	DRIVER	f25f09798b51603bde602ded37fea826fc8b7f962fc3b0...	CHICA
5	O877654	DRIVER	e9146986f4b0884d00ff3a54da5249263b4b36c15d01ce...	CHICA
8	O879685	DRIVER	cb7e25c6e63094307e5b999a8a5ec6f781a54781aeaef3...	MELRO PA
10	O880033	DRIVER	2bdb105a834f9917ee9be1937302d775525192c91e7398...	CARBONDA

5 rows × 45 columns

In [34]: # checking for null values

```
chicagocrush.info()
   # STREET_NAME      529191 non-null object
  28 MOST_SEVERE_INJURY      529191 non-null object
  29 INJURIES_TOTAL      529191 non-null float64
  30 CRASH_HOUR      529191 non-null float64
  31 CRASH_DAY_OF_WEEK      529191 non-null float64
  32 CRASH_MONTH      529191 non-null float64
  33 LOCATION      529191 non-null object
  34 UNIT_NO      529191 non-null float64
  35 UNIT_TYPE      529191 non-null object
  36 VEHICLE_ID      529191 non-null float64
  37 MAKE      529191 non-null object
  38 MODEL      529191 non-null object
  39 VEHICLE_YEAR      529191 non-null float64
  40 VEHICLE_DEFECT      529191 non-null object
  41 VEHICLE_TYPE      529191 non-null object
  42 VEHICLE_USE      529191 non-null object
  43 MANEUVER      529191 non-null object
  44 OCCUPANT_CNT      529191 non-null float64
dtypes: float64(10), object(35)
memory usage: 185.7+ MB
```

In [35]: # dropping irrelevant columns in the merged dataset

```
columns_drop = ['PERSON_ID', 'PERSON_TYPE', 'DRIVERS_LICENSE_STATE', 'DRIVERS_LICEN
chicagocrush1 = chicagocrush.drop(columns=columns_drop)
chicagocrush1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 529191 entries, 0 to 1553242
Data columns (total 39 columns):
 #   Column           Non-Null Count  Dtype  
---  -- 
 0   CRASH_RECORD_ID    529191 non-null object 
 1   CITY              529191 non-null object 
 2   STATE             529191 non-null object 
 3   SEX               529191 non-null object 
 4   AGE               529191 non-null float64
 5   AIRBAG_DEPLOYED   529191 non-null object 
 6   EJECTION          529191 non-null object 
 7   INJURY_CLASSIFICATION 529191 non-null object 
 8   DRIVER_ACTION     529191 non-null object 
 9   DRIVER_VISION     529191 non-null object 
 10  PHYSICAL_CONDITION 529191 non-null object 
 11  CRASH_DATE        529191 non-null object 
 12  POSTED_SPEED_LIMIT 529191 non-null float64
 13  TRAFFIC_CONTROL_DEVICE 529191 non-null object 
 14  DRIVER_CONDITION   529191 non-null object 
```

```
In [36]: chicagocrush1['CRASH_DATE'] = pd.to_datetime(chicagocrush1['CRASH_DATE'])

# Filter the data for the year 2023
Chicagocrush2023 = chicagocrush1[chicagocrush1['CRASH_DATE'].dt.year == 2023]

# Display the first few rows of the filtered DataFrame
Chicagocrush2023.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 68767 entries, 324 to 1553242
Data columns (total 39 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CRASH_RECORD_ID    68767 non-null   object 
 1   CITY               68767 non-null   object 
 2   STATE              68767 non-null   object 
 3   SEX                68767 non-null   object 
 4   AGE                68767 non-null   float64
 5   AIRBAG_DEPLOYED   68767 non-null   object 
 6   EJECTION            68767 non-null   object 
 7   INJURY_CLASSIFICATION 68767 non-null   object 
 8   DRIVER_ACTION      68767 non-null   object 
 9   DRIVER_VISION      68767 non-null   object 
 10  PHYSICAL_CONDITION 68767 non-null   object 
 11  CRASH_DATE         68767 non-null   datetime64[ns]
 12  POSTED_SPEED_LIMIT 68767 non-null   float64
 13  TRAFFIC_CONTROL_DEVICE 68767 non-null   object 
 14  DEVICE_CONDITION   68767 non-null   object 
 15  WEATHER_CONDITION  68767 non-null   object 
 16  LIGHTING_CONDITION 68767 non-null   object 
 17  ROAD_DEFECT        68767 non-null   object 
 18  CRASH_TYPE          68767 non-null   object 
 19  DAMAGE              68767 non-null   object 
 20  PRIM_CONTRIBUTORY_CAUSE 68767 non-null   object 
 21  SEC_CONTRIBUTORY_CAUSE 68767 non-null   object 
 22  STREET_NAME         68767 non-null   object 
 23  MOST_SEVERE_INJURY 68767 non-null   object 
 24  INJURIES_TOTAL     68767 non-null   float64
 25  CRASH_HOUR          68767 non-null   float64
 26  CRASH_DAY_OF_WEEK   68767 non-null   float64
 27  CRASH_MONTH         68767 non-null   float64
 28  LOCATION             68767 non-null   object 
 29  UNIT_NO              68767 non-null   float64
 30  UNIT_TYPE            68767 non-null   object 
 31  VEHICLE_ID          68767 non-null   float64
 32  MAKE                68767 non-null   object 
 33  MODEL                68767 non-null   object 
 34  VEHICLE_YEAR         68767 non-null   float64
 35  VEHICLE_DEFECT       68767 non-null   object 
 36  VEHICLE_TYPE          68767 non-null   object 
 37  VEHICLE_USE           68767 non-null   object 
 38  MANEUVER             68767 non-null   object 

dtypes: datetime64[ns](1), float64(9), object(29)
memory usage: 21.0+ MB
```

In [37]:

```
# Encode categorical variables
label_encoder = LabelEncoder()

Chicagocrush2023.loc[:, 'INJURY_CLASSIFICATION'] = label_encoder.fit_transform(Chicagocrush2023['INJURY_CLASSIFICATION'])
Chicagocrush2023.loc[:, 'DRIVER_VISION'] = label_encoder.fit_transform(Chicagocrush2023['DRIVER_VISION'])
Chicagocrush2023.loc[:, 'DRIVER_ACTION'] = label_encoder.fit_transform(Chicagocrush2023['DRIVER_ACTION'])
Chicagocrush2023.loc[:, 'PHYSICAL_CONDITION'] = label_encoder.fit_transform(Chicagocrush2023['PHYSICAL_CONDITION'])
Chicagocrush2023.loc[:, 'LIGHTING_CONDITION'] = label_encoder.fit_transform(Chicagocrush2023['LIGHTING_CONDITION'])
Chicagocrush2023.loc[:, 'AIRBAG_DEPLOYED'] = label_encoder.fit_transform(Chicagocrush2023['AIRBAG_DEPLOYED'])
Chicagocrush2023.loc[:, 'MAKE'] = label_encoder.fit_transform(Chicagocrush2023['MAKE'])
Chicagocrush2023.loc[:, 'CITY'] = label_encoder.fit_transform(Chicagocrush2023['CITY'])
Chicagocrush2023.loc[:, 'SEX'] = label_encoder.fit_transform(Chicagocrush2023['SEX'])
Chicagocrush2023.loc[:, 'TRAFFIC_CONTROL_DEVICE'] = label_encoder.fit_transform(Chicagocrush2023['TRAFFIC_CONTROL_DEVICE'])
Chicagocrush2023.loc[:, 'ROAD_DEFECT'] = label_encoder.fit_transform(Chicagocrush2023['ROAD_DEFECT'])
Chicagocrush2023.loc[:, 'VEHICLE_USE'] = label_encoder.fit_transform(Chicagocrush2023['VEHICLE_USE'])
```

In [38]: Chicagocrush2023.head()

Out[38]:

		CRASH_RECORD_ID	CITY	STATE	SEX	AGE	AIRBAG_DE
324		a073a023f6713d0957a093f7e550175fd8699ec037c9d5...	222	CT	1	32.0	
1389		c026bf2731fafc1ce64b302a52f58c9110627fca47541e...	384	IL	1	24.0	
27431		b51f3812261afa62974edad0c600168daac53c678adea3...	384	IL	0	36.0	
31253		41c564000224de0aa0bd4a0af0a8985a4538a2f89512e0...	1140	IL	1	24.0	
40908		0fc7fd50aac51f2a839c01ce53180d43f87803c598dd6b...	384	IL	0	40.0	

5 rows × 39 columns

```
In [39]: # columns in the cleaned merged dataset
threshold = 0 # Set the threshold

# Create a new binary column indicating whether there are injuries (1) or not (0)
Chicagocrush2023.loc[:, 'INJURIES'] = (Chicagocrush2023['INJURIES_TOTAL'] > threshold).astype(int)

# Display the first few rows of the DataFrame with the new binary column
print(Chicagocrush2023[['INJURIES_TOTAL', 'INJURIES']].head())
```

	INJURIES_TOTAL	INJURIES
324	0.0	0
1389	0.0	0
27431	0.0	0
31253	0.0	0
40908	0.0	0

```
C:\Users\Home\AppData\Local\Temp\ipykernel_1508\2471359133.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Chicagocrush2023.loc[:, 'INJURIES'] = (Chicagocrush2023['INJURIES_TOTAL'] > threshold).astype(int)
```

Data Exploration

```
In [40]: # group different types of traffic control devices into a smaller set of categories
traffic_control_map = {
    'NO CONTROLS': 'No Controls',
    'TRAFFIC SIGNAL': 'Traffic Signal',
    'STOP SIGN/FLASHER': 'Stop Sign/Flasher',
    'UNKNOWN': 'Unknown/Other',
    'OTHER': 'Unknown/Other',
    'LANE USE MARKING': 'Traffic Signal',
    'YIELD': 'Unknown/Other',
    'OTHER REG. SIGN': 'Traffic Signal',
    'OTHER WARNING SIGN': 'Traffic Signal',
    'RAILROAD CROSSING GATE': 'Traffic Signal',
    'PEDESTRIAN CROSSING SIGN': 'Traffic Signal',
    'DELINEATORS': 'Traffic Signal',
    'FLASHING CONTROL SIGNAL': 'Traffic Signal',
    'POLICE/FLAGMAN': 'Traffic Signal',
    'SCHOOL ZONE': 'Traffic Signal',
    'OTHER RAILROAD CROSSING': 'Traffic Signal',
    'RR CROSSING SIGN': 'Traffic Signal',
    'NO PASSING': 'Traffic Signal',
    'BICYCLE CROSSING SIGN': 'Traffic Signal'
}

# Remap values in the 'TRAFFIC_CONTROL_DEVICE' column
Chicagocrush2023.loc[:, 'TRAFFIC_CONTROL_DEVICE'] = chicagocrush1['TRAFFIC_CONTROL_DEVICE'].map(traffic_control_map)

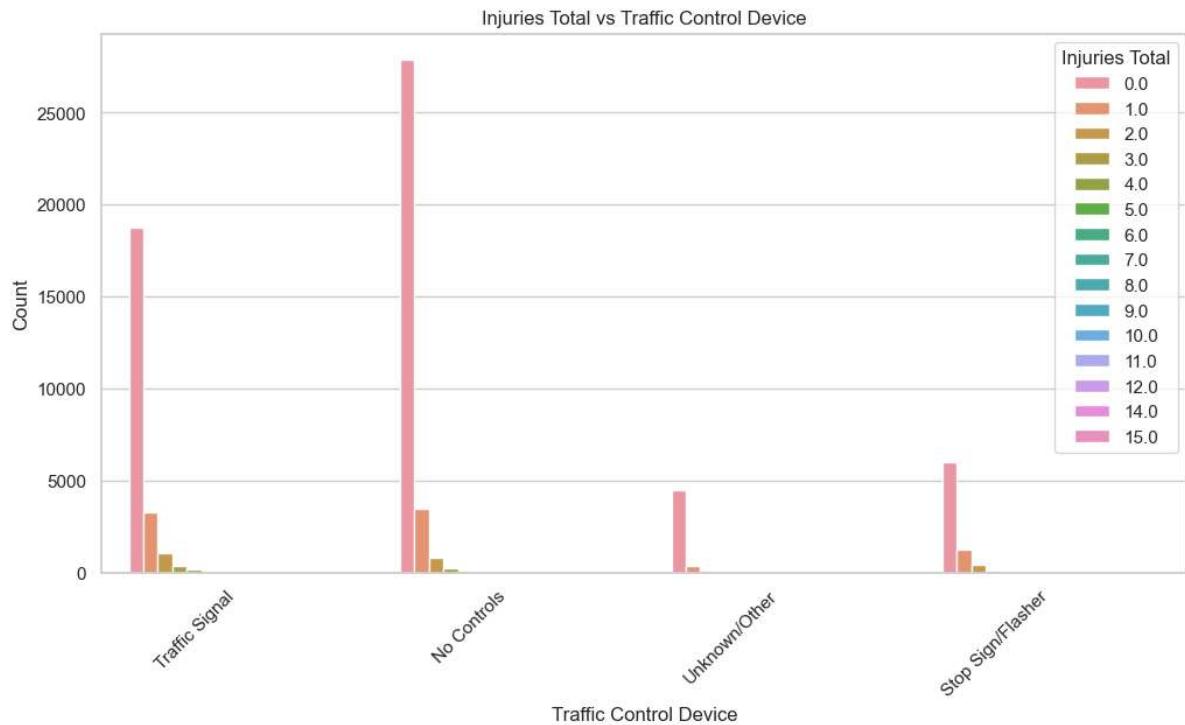
# Check the updated value counts
print(Chicagocrush2023['TRAFFIC_CONTROL_DEVICE'].value_counts())
```

```
TRAFFIC_CONTROL_DEVICE
No Controls      32454
Traffic Signal   23561
Stop Sign/Flasher 7832
Unknown/Other     4920
Name: count, dtype: int64
```

In [41]: # Create a countplot to visualize the distribution of injuries for each traffic control device

```
sns.set(style="whitegrid")

plt.figure(figsize=(12, 6))
sns.countplot(x='TRAFFIC_CONTROL_DEVICE', hue='INJURIES_TOTAL', data=ChicagoCrash)
plt.title('Injuries Total vs Traffic Control Device')
plt.xlabel('Traffic Control Device')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better visibility
plt.legend(title='Injuries Total', loc='upper right')
plt.show()
```



```
In [42]: # Defining a dictionary map for remapping weather conditions to plot against it
weather_condition_map = {
    'CLEAR': 'CLEAR',
    'RAIN': 'RAIN',
    'SNOW': 'SNOW',
    'CLOUDY/OVERCAST': 'CLOUDY',
    'UNKNOWN': 'CLEAR',
    'OTHER': 'CLEAR',
    'FREEZING RAIN/DRIZZLE': 'RAIN',
    'FOG/SMOKE/HAZE': 'FOG',
    'SLEET/HAIL': 'FOG',
    'BLOWING SNOW': 'FOG',
    'SEVERE CROSS WIND GATE': 'FOG',
    'BLOWING SAND, SOIL, DIRT': 'FOG'
}

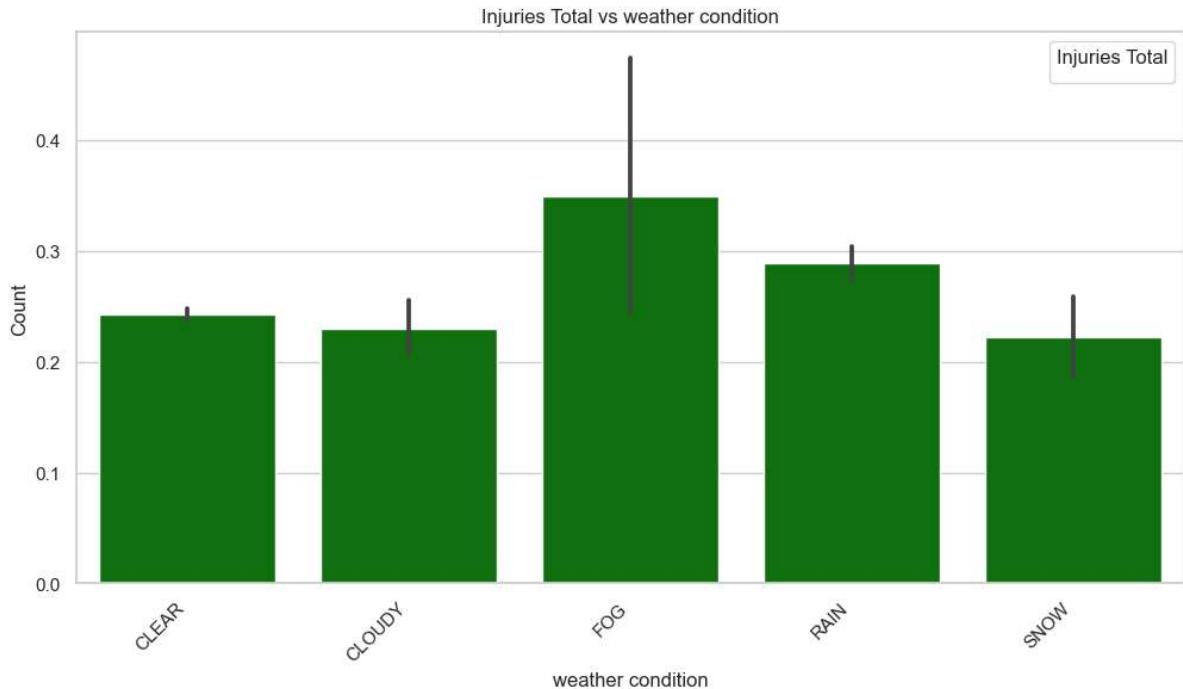
# Remap values in the 'WEATHER_CONDITION' column using the defined map
Chicagocrush2023.loc[:, 'WEATHER_CONDITION'] = Chicagocrush2023['WEATHER_CONDITION'].map(weather_condition_map)

# Check the updated value counts
weather_condition_counts = Chicagocrush2023['WEATHER_CONDITION'].value_counts()
print(weather_condition_counts)
```

```
WEATHER_CONDITION
CLEAR      57350
RAIN       7444
CLOUDY     2301
SNOW       1406
FOG        266
Name: count, dtype: int64
```

```
In [43]: plt.figure(figsize=(12, 6))
sns.barplot(x='WEATHER_CONDITION', y='INJURIES_TOTAL', data=Chicagocrush2023, )
plt.title('Injuries Total vs weather condition')
plt.xlabel('weather condition')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better visibility
plt.legend(title='Injuries Total', loc='upper right')
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



```
In [44]: # FEATURE ENGINEERING
# Binning age
bins = [0, 25, 50, 100]
labels = ['Young', 'Adult', 'Senior']
Chicagocrush2023.loc[:, 'AGE_GROUP'] = pd.cut(Chicagocrush2023['AGE'], bins=bins)
```

C:\Users\Home\AppData\Local\Temp\ipykernel_1508\3919465822.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

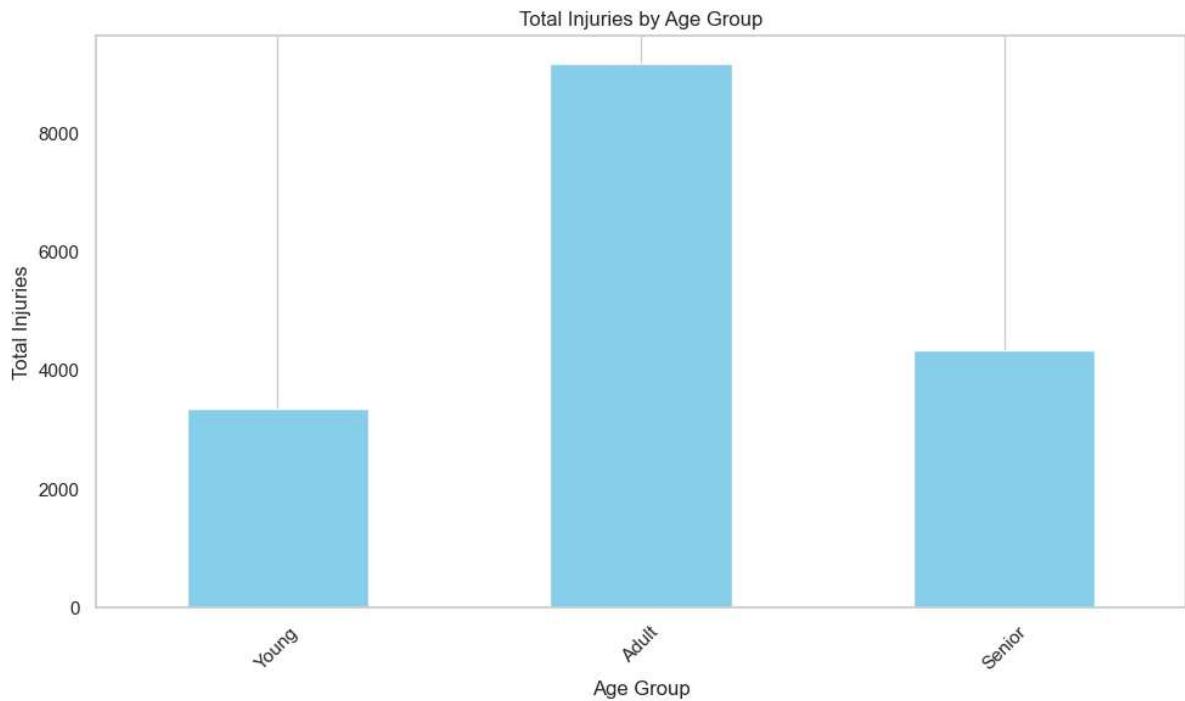
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Chicagocrush2023.loc[:, 'AGE_GROUP'] = pd.cut(Chicagocrush2023['AGE'], bins=bins, labels=labels)
```

In [45]:

```
# Assuming you have a DataFrame named 'data' with 'AGE_GROUP' and 'TOTAL_INJURIES' columns
# Group the data by 'AGE_GROUP' and calculate the total injuries for each group
agegroup_total_injuries = Chicagocrush2023.groupby('AGE_GROUP')[['INJURIES_TOTAL']].sum()

# Plot the bar chart
plt.figure(figsize=(10, 6))
agegroup_total_injuries.plot(kind='bar', color='skyblue')
plt.title('Total Injuries by Age Group')
plt.xlabel('Age Group')
plt.ylabel('Total Injuries')
plt.xticks(rotation=45) # Rotate x-axis labels for better visibility
plt.grid(axis='y') # Add gridlines along y-axis
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```



In [46]: # Time period encoding

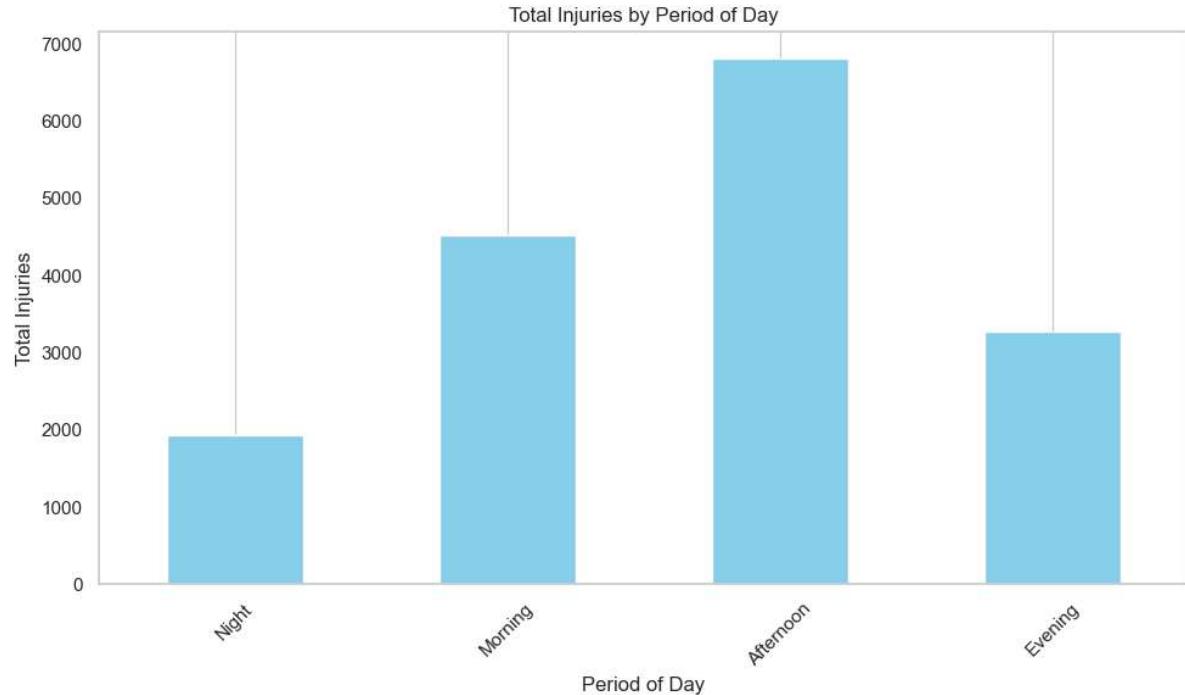
```
Chicagocrush2023.loc[:, 'PERIOD_OF_DAY'] = pd.cut(Chicagocrush2023['CRASH_HOUR'],
# Group the data by 'PERIOD_OF_DAY' and calculate the total injuries for each group
period_of_day_total_injuries = Chicagocrush2023.groupby('PERIOD_OF_DAY')[['INJURIES']].sum()

# Plot the bar chart
plt.figure(figsize=(10, 6))
period_of_day_total_injuries.plot(kind='bar', color='skyblue')
plt.title('Total Injuries by Period of Day')
plt.xlabel('Period of Day')
plt.ylabel('Total Injuries')
plt.xticks(rotation=45) # Rotate x-axis labels for better visibility
plt.grid(axis='y') # Add gridlines along y-axis
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```

C:\Users\Home\AppData\Local\Temp\ipykernel_1508\3928030688.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

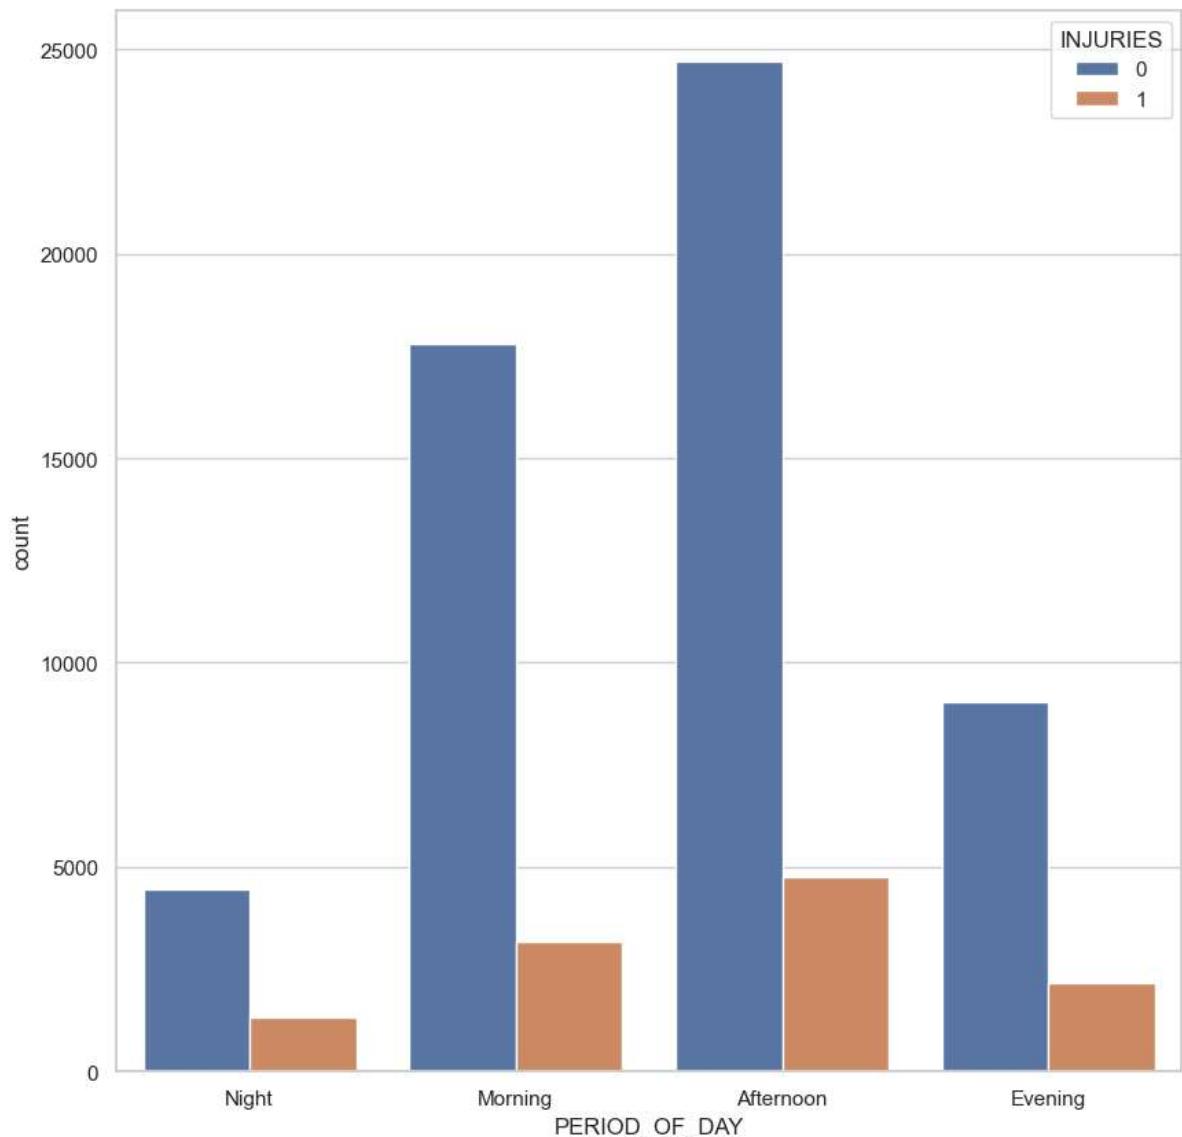
```
Chicagocrush2023.loc[:, 'PERIOD_OF_DAY'] = pd.cut(Chicagocrush2023['CRASH_HOUR'],
bins=[0, 6, 12, 18, 24], labels=['Night', 'Morning', 'Afternoon', 'Evening'])
```



In [47]: # representing period of day in relation to injuries

```
plt.figure(figsize=(10,10))
sns.countplot(x="PERIOD_OF_DAY", hue="INJURIES", data=Chicagocrush2023)
```

Out[47]: <Axes: xlabel='PERIOD_OF_DAY', ylabel='count'>



```
In [48]: # Convert 'CRASH_DATE' to datetime format
Chicagocrush2023.loc[:, 'CRASH_DATE'] = pd.to_datetime(Chicagocrush2023['CRASH_DATE'])

# Extract just the year
Chicagocrush2023.loc[:, 'YEAR'] = Chicagocrush2023['CRASH_DATE'].dt.year
# Calculate the age of the vehicle at the time of the accident
Chicagocrush2023.loc[:, 'VEHICLE AGE'] = Chicagocrush2023['YEAR'] - Chicagocrush2023['BIRTH_YEAR']

# Binning vehicle age
bins = [0, 5, 10, float('inf')] # Example bins: 0-5 years, 5-10 years, >10 years
labels = ['New', 'Moderately Aged', 'Old']
Chicagocrush2023.loc[:, 'VEHICLE AGE GROUP'] = pd.cut(Chicagocrush2023['VEHICLE AGE'], bins, labels=labels)
# Group the data by 'VEHICLE AGE GROUP' and calculate the total injuries for each group
vehicle_age_group_total_injuries = Chicagocrush2023.groupby('VEHICLE AGE GROUP').sum()

# Plot the bar chart
plt.figure(figsize=(10, 6))
vehicle_age_group_total_injuries.plot(kind='bar', color='skyblue')
plt.title('Total Injuries by vehicle age group ')
plt.xlabel('Vehicle age group')
plt.ylabel('Total Injuries')
plt.xticks(rotation=45) # Rotate x-axis labels for better visibility
plt.grid(axis='y') # Add gridlines along y-axis
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```

```
C:\Users\Home\AppData\Local\Temp\ipykernel_1508\1816627131.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Chicagocrush2023.loc[:, 'YEAR'] = Chicagocrush2023['CRASH_DATE'].dt.year
```

```
C:\Users\Home\AppData\Local\Temp\ipykernel_1508\1816627131.py:7: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Chicagocrush2023.loc[:, 'VEHICLE AGE'] = Chicagocrush2023['YEAR'] - Chicagocrush2023['VEHICLE_YEAR']
```

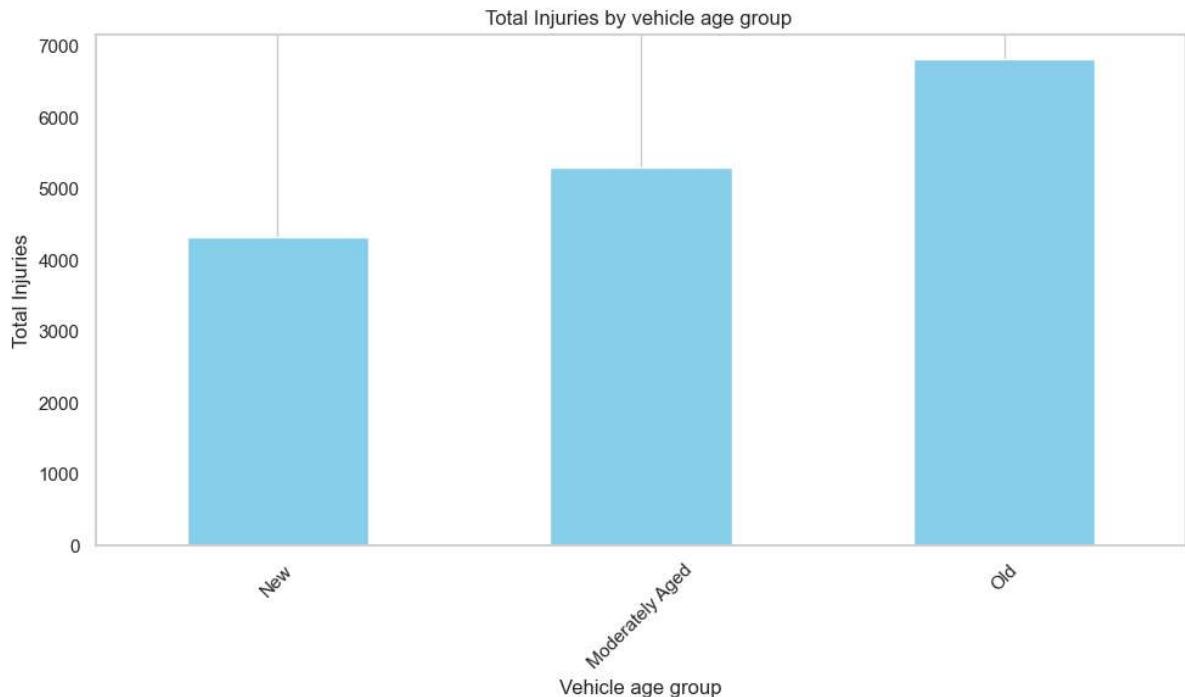
```
C:\Users\Home\AppData\Local\Temp\ipykernel_1508\1816627131.py:12: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Chicagocrush2023.loc[:, 'VEHICLE AGE GROUP'] = pd.cut(Chicagocrush2023['VEHICLE AGE'], bins=bins, labels=labels)
```



```
In [49]: # preview the columns in the dataset
```

```
Chicagocrush2023.columns
```

```
Out[49]: Index(['CRASH_RECORD_ID', 'CITY', 'STATE', 'SEX', 'AGE', 'AIRBAG_DEPLOYED',  
   'EJECTION', 'INJURY_CLASSIFICATION', 'DRIVER_ACTION', 'DRIVER_VISION',  
   'PHYSICAL_CONDITION', 'CRASH_DATE', 'POSTED_SPEED_LIMIT',  
   'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION',  
   'LIGHTING_CONDITION', 'ROAD_DEFECT', 'CRASH_TYPE', 'DAMAGE',  
   'PRIM_CONTRIBUTORY_CAUSE', 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NAME',  
   'MOST_SEVERE_INJURY', 'INJURIES_TOTAL', 'CRASH_HOUR',  
   'CRASH_DAY_OF_WEEK', 'CRASH_MONTH', 'LOCATION', 'UNIT_NO', 'UNIT_TYP  
E',  
   'VEHICLE_ID', 'MAKE', 'MODEL', 'VEHICLE_YEAR', 'VEHICLE_DEFECT',  
   'VEHICLE_TYPE', 'VEHICLE_USE', 'MANEUVER', 'INJURIES', 'AGE_GROUP',  
   'PERIOD_OF_DAY', 'YEAR', 'VEHICLE_AGE', 'VEHICLE_AGE_GROUP'],  
  dtype='object')
```

```
In [50]: # dropping columns
```

```
columns_drop = ['DEVICE_CONDITION', 'VEHICLE_DEFECT', 'MODEL']
Chicagocrush2023 = Chicagocrush2023.drop(columns=columns_drop)
Chicagocrush2023.info()

<class 'pandas.core.frame.DataFrame'>
Index: 68767 entries, 324 to 1553242
Data columns (total 42 columns):
 #   Column           Non-Null Count Dtype  
 ---  -- 
 0   CRASH_RECORD_ID    68767 non-null  object  
 1   CITY               68767 non-null  object  
 2   STATE              68767 non-null  object  
 3   SEX                68767 non-null  object  
 4   AGE                68767 non-null  float64 
 5   AIRBAG_DEPLOYED   68767 non-null  object  
 6   EJECTION            68767 non-null  object  
 7   INJURY_CLASSIFICATION 68767 non-null  object  
 8   DRIVER_ACTION      68767 non-null  object  
 9   DRIVER_VISION       68767 non-null  object  
 10  PHYSICAL_CONDITION 68767 non-null  object  
 11  CRASH_DATE         68767 non-null  datetime64[ns] 
 12  POSTED_SPEED_LIMIT 68767 non-null  float64 
 13  TRAFFIC_CONTROL_DEVICE 68767 non-null  object  
 14  WEATHER_CONDITION   68767 non-null  object  
 15  LIGHTING_CONDITION  68767 non-null  object  
 16  ROAD_DEFECT          68767 non-null  object  
 17  CRASH_TYPE           68767 non-null  object  
 18  DAMAGE               68767 non-null  object  
 19  PRIM_CONTRIBUTORY_CAUSE 68767 non-null  object  
 20  SEC_CONTRIBUTORY_CAUSE 68767 non-null  object  
 21  STREET_NAME          68767 non-null  object  
 22  MOST_SEVERE_INJURY    68767 non-null  object  
 23  INJURIES_TOTAL        68767 non-null  float64 
 24  CRASH_HOUR            68767 non-null  float64 
 25  CRASH_DAY_OF_WEEK     68767 non-null  float64 
 26  CRASH_MONTH            68767 non-null  float64 
 27  LOCATION              68767 non-null  object  
 28  UNIT_NO               68767 non-null  float64 
 29  UNIT_TYPE              68767 non-null  object  
 30  VEHICLE_ID             68767 non-null  float64 
 31  MAKE                  68767 non-null  object  
 32  VEHICLE_YEAR            68767 non-null  float64 
 33  VEHICLE_TYPE            68767 non-null  object  
 34  VEHICLE_USE              68767 non-null  object  
 35  MANEUVER                68767 non-null  object  
 36  INJURIES                 68767 non-null  int32  
 37  AGE_GROUP                68223 non-null  category 
 38  PERIOD_OF_DAY            67439 non-null  category 
 39  YEAR                     68767 non-null  int32  
 40  VEHICLE_AGE                68767 non-null  float64 
 41  VEHICLE_AGE_GROUP        66247 non-null  category 

dtypes: category(3), datetime64[ns](1), float64(10), int32(2), object(26)
memory usage: 20.7+ MB
```

In [51]: # previewing the data columns and type

```
ChicagoCrash2023.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 68767 entries, 324 to 1553242
Data columns (total 42 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CRASH_RECORD_ID    68767 non-null   object 
 1   CITY               68767 non-null   object 
 2   STATE              68767 non-null   object 
 3   SEX                68767 non-null   object 
 4   AGE                68767 non-null   float64
 5   AIRBAG_DEPLOYED   68767 non-null   object 
 6   EJECTION           68767 non-null   object 
 7   INJURY_CLASSIFICATION 68767 non-null   object 
 8   DRIVER_ACTION      68767 non-null   object 
 9   DRIVER_VISION      68767 non-null   object 
 10  PHYSICAL_CONDITION 68767 non-null   object 
 11  CRASH_DATE         68767 non-null   datetime64[ns]
 12  POSTED_SPEED_LIMIT 68767 non-null   float64
 13  TRAFFIC_CONTROL_DEVICE 68767 non-null   object 
 14  WEATHER_CONDITION   68767 non-null   object 
 15  LIGHTING_CONDITION  68767 non-null   object 
 16  ROAD_DEFECT         68767 non-null   object 
 17  CRASH_TYPE          68767 non-null   object 
 18  DAMAGE              68767 non-null   object 
 19  PRIM_CONTRIBUTORY_CAUSE 68767 non-null   object 
 20  SEC_CONTRIBUTORY_CAUSE 68767 non-null   object 
 21  STREET_NAME          68767 non-null   object 
 22  MOST_SEVERE_INJURY   68767 non-null   object 
 23  INJURIES_TOTAL       68767 non-null   float64
 24  CRASH_HOUR           68767 non-null   float64
 25  CRASH_DAY_OF_WEEK    68767 non-null   float64
 26  CRASH_MONTH          68767 non-null   float64
 27  LOCATION             68767 non-null   object 
 28  UNIT_NO              68767 non-null   float64
 29  UNIT_TYPE            68767 non-null   object 
 30  VEHICLE_ID           68767 non-null   float64
 31  MAKE                68767 non-null   object 
 32  VEHICLE_YEAR          68767 non-null   float64
 33  VEHICLE_TYPE          68767 non-null   object 
 34  VEHICLE_USE           68767 non-null   object 
 35  MANEUVER             68767 non-null   object 
 36  INJURIES             68767 non-null   int32  
 37  AGE_GROUP             68223 non-null   category
 38  PERIOD_OF_DAY         67439 non-null   category
 39  YEAR                 68767 non-null   int32  
 40  VEHICLE_AGE           68767 non-null   float64
 41  VEHICLE_AGE_GROUP     66247 non-null   category
dtypes: category(3), datetime64[ns](1), float64(10), int32(2), object(26)
memory usage: 20.7+ MB
```

In [52]: # dropping irrelevant columns

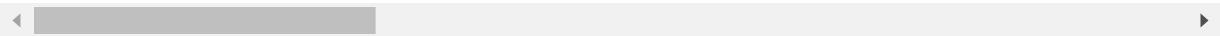
```
columns_drop = ['STATE', 'EJECTION', 'INJURY_CLASSIFICATION', 'CRASH_TYPE', 'DAMAGE',  
                 'SEC_CONTRIBUTORY_CAUSE', 'STREET_NAME', 'MOST_SEVERE_INJURY', '  
                 'UNIT_NO', 'UNIT_TYPE', 'VEHICLE_ID', 'MAKE', 'MANEUVER', 'INJURIES',  
                 'CRASH_DATE', 'CRASH_RECORD_ID', 'VEHICLE_TYPE', 'VEHICLE_YEAR', '  
  
Chicagocrush2023 = Chicagocrush2023.drop(columns=columns_drop)
```

In [53]: # data preview

```
Chicagocrush2023.head()
```

Out[53]:

	CITY	SEX	AIRBAG_DEPLOYED	DRIVER_ACTION	DRIVER_VISION	PHYSICAL_CONDITION	
324	222	1		5	18	12	{
1389	384	1		5	13	12	{
27431	384	0		3	7	12	{
31253	1140	1		5	12	7	{
40908	384	0		5	18	7	{



Train test split

In [54]:

```
# Define features (X) and target variable (y)
X = Chicagocrush2023.drop(columns=['INJURIES']) # Features
y = Chicagocrush2023['INJURIES'] # Target variable

# Split the data into training and temporary set (70% training, 30% temporary)
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42)

# Split the temporary set into validation and testing sets (50% validation, 50% testing)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

# Display the shapes of the resulting sets
print("Training set:", X_train.shape, y_train.shape)
print("Validation set:", X_val.shape, y_val.shape)
print("Testing set:", X_test.shape, y_test.shape)
```

Training set: (48136, 17) (48136,)
 Validation set: (10315, 17) (10315,)
 Testing set: (10316, 17) (10316,)

In [55]: num_columns = Chicagocrush2023.drop(columns='INJURIES').select_dtypes('number')

Out[55]: ['POSTED_SPEED_LIMIT', 'VEHICLE_AGE']

In [56]: cat_cols = Chicagocrush2023.drop(columns='INJURIES').select_dtypes('object').columns

Out[56]: ['CITY',
 'SEX',
 'AIRBAG_DEPLOYED',
 'DRIVER_ACTION',
 'DRIVER_VISION',
 'PHYSICAL_CONDITION',
 'TRAFFIC_CONTROL_DEVICE',
 'WEATHER_CONDITION',
 'LIGHTING_CONDITION',
 'ROAD_DEFECT',
 'LOCATION',
 'VEHICLE_USE']

In [57]:

```
numerical_transformer = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='mean')),  
    ('scaler', StandardScaler())  
)  
  
categorical_transformer = Pipeline(steps=[  
    ('imputer', SimpleImputer(strategy='most_frequent')),  
    ('onehot', OneHotEncoder(handle_unknown='ignore'))  
)  
  
# combine both pipelines into one using columntransformer  
preprocessor = ColumnTransformer(  
    transformers=[  
        ('num', numerical_transformer, num_columns),  
        ('cat', categorical_transformer, cat_cols)  
    ])  
  
  
# preprocessing X_train and X_test  
X_train_transformed = preprocessor.fit_transform(X_train)  
X_test_transformed = preprocessor.transform(X_test)  
# Fit and transform the data  
# X_transformed = preprocessor.fit_transform(Chicagocrush2023)
```

In [58]: X_test_transformed

Out[58]: <10316x36832 sparse matrix of type '<class 'numpy.float64'>'
with 138145 stored elements in Compressed Sparse Row format>

```
In [59]: # accessing categorical columns from pipeline then converting to dataframe
slice_pipe = preprocessor.named_transformers_['cat']
cat_features = slice_pipe.named_steps['onehot'].get_feature_names_out(cat_cols)
X_train_transformed = pd.DataFrame(X_train_transformed,columns=[*num_columns, *cat_features])
X_train_transformed
```

```

-----
```

ValueError Traceback (most recent call last)

Cell In[59], line 4

```

    2 slice_pipe = preprocessor.named_transformers_['cat']
    3 cat_features = slice_pipe.named_steps['onehot'].get_feature_names_out
(cat_cols)
--> 4 X_train_transformed = pd.DataFrame(X_train_transformed,columns=[*num_
columns,*cat_features])
    5 X_train_transformed

```

File ~\anaconda3\Lib\site-packages\pandas\core\frame.py:798, in DataFrame.__i
nit__(self, data, index, columns, dtype, copy)

```

790     mgr = arrays_to_mgr(
791         arrays,
792         columns,
(...),
795         typ=manager,
796     )
797     else:
--> 798     mgr = ndarray_to_mgr(
799         data,
800         index,
801         columns,
802         dtype=dtype,
803         copy=copy,
804         typ=manager,
805     )
806     else:
807     mgr = dict_to_mgr(
808         {},
809         index,
(...),
812         typ=manager,
813     )

```

File ~\anaconda3\Lib\site-packages\pandas\core\internals\construction.py:337,
in ndarray_to_mgr(values, index, columns, dtype, copy, typ)

```

332 # _prep_ndarraylike ensures that values.ndim == 2 at this point
333 index, columns = _get_axes(
334     values.shape[0], values.shape[1], index=index, columns=columns
335 )
--> 337 _check_values_indices_shape_match(values, index, columns)
339 if typ == "array":
340     if issubclass(values.dtype.type, str):

```

File ~\anaconda3\Lib\site-packages\pandas\core\internals\construction.py:408,
in _check_values_indices_shape_match(values, index, columns)

```

406 passed = values.shape
407 implied = (len(index), len(columns))
--> 408 raise ValueError(f"Shape of passed values is {passed}, indices imply
{implied}")

```

ValueError: Shape of passed values is (48136, 1), indices imply (48136, 3683
2)

In []:

```
# Checking for imbalance visually
plt.figure(figsize=(6, 4))
sns.countplot(x=y_train)
plt.title('Class Distribution')
plt.show()
```

In []:

```
# checking for imbalance
class_distribution = y_train.value_counts(normalize=True)
print(class_distribution)
```

In []:

```
# observation there is high imbalance and to solve we will use resampling(SMOTE)

# from imblearn.over_sampling import SMOTE
# smote = SMOTE(sampling_strategy='auto', random_state=42)
# X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

LOGISTIC REGRESSION MODEL

In []:

```
# instantiate LogisticRegression
log_reg = LogisticRegression()
```

In []:

```
# Fit the model
log_reg.fit(X_train, y_train)
# predict
# y_pred = log_reg.predict(X_test)
```

KNN MODEL

In []:

```
knn_clf = KNeighborsClassifier()
```

In []:

```
# fit the model
knn_clf.fit(X_train_transformed, y_train)

# predict
y_pred = knn_clf.predict(X_test_transformed)
```

```
In [ ]: # Compute accuracy
accuracy = accuracy_score(y_test, y_pred)

# Compute precision, recall, and F1-score
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

# Compute confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Print classification report
class_report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)
```

Decision trees

```
In [ ]: # instantiate DecisionTreeClassifier
tree_clf = DecisionTreeClassifier()
```

```
In [ ]: # fit the model
tree_clf.fit(X_train_transformed, y_train)

# predict
y_pred = tree_clf.predict(X_test_transformed)
```

```
In [ ]: # Compute accuracy
accuracy = accuracy_score(y_test, y_pred)

# Compute precision, recall, and F1-score
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

# Compute confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Print classification report
class_report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)
```

Intrepretation

KNN and Decision trees returned an accuracy rate of 83%

Recommendations

The following are our recommendations ;

- The city to put restrictions on vehicles allowed in roads to be below ten years
- The city lowers the speed limit during afternoon/rush hour or more patrol in the 30-40 mph zones.
- The city should increase the traffic control devices to reduce in accidents occurrences

Conclusion

- Most of the injuries recorded were caused by drivers between the age of 25 to 50 years falling into the adult category
- Based on the data, it shows that most accidents occur in the afternoon/rush hour.
- It also shows that most accidents occur in speed limit zones labeled between 30-40 mph.
- Most of the vehicles involved in the accidents fell into the old category that is they were in the road for a period of more than ten years
- Most of the accidents occurred where there were no traffic signals but less injuries were observed

